

Generalization

Reinforcement Learning

Computer Engineering Department
Sharif University of Technology



States

are continuous

or

there are infinitely


many of them.

Mohammad Hossein Rohban, Ph.D.

Spring 2024

Courtesy: Some slides are adopted from CS 285 Berkeley, and CS 234 Stanford, and Pieter Abbeel's compact series on RL.

feature
 $\Phi(s)$ Representation



$\rightarrow \pi = \epsilon\text{-greedy } \hat{Q}_w(s, a)$ Replay Buffer
 $\rightarrow s_0, a_0, r_0, s_1, a_1, r_1, \dots$
training sample

Function Approximation

MC TD for the network
 $(\underbrace{\Phi(s_0)}_{\text{state}}, \underbrace{G_0}_{\text{action}}) \rightarrow (\Phi(s_0, a_0), r_0 + \gamma \max_a \hat{Q}_w(s_1, a))$
 $(\Phi(s_1, a_1), G_1) \rightarrow (\Phi(s_1, a_1), r_0 + \gamma \hat{Q}_w(s_1, a_1))$
 \vdots
 $\hat{Q}_w(s, a)$

Function approximation and deep RL

- The **policy**, **value function**, **model**, and **agent state update** are all functions
- We want to learn these from experience (data).
- If there are too many states, we need to approximate.
- This is often called **deep reinforcement learning**, when using **neural networks** to represent these functions.

Large-Scale Reinforcement Learning

- Reinforcement learning can be used to solve **large** problems, e.g.
 - Backgammon: 10^{20} states
 - Go: 10^{170} states
 - Helicopter: continuous state space
 - Robots: real world
- How can we apply our methods for **prediction** and **control**?

Value Function Approximation

Value Function Approximation

- So far we mostly considered **lookup tables**
 - Every state s has an entry $v(s)$
 - Or every state-action pair s, a has an entry $q(s, a)$
- Problem with large MDPs:
 - There are too many states and/or actions to store in memory
 - It is too slow to learn the value of each state individually
 - Individual environment states are often **not fully observable**

Value Function Approximation

- Solution for large MDPs:
 - Estimate value function with **function approximation**

$$\begin{array}{ll} v_{\mathbf{w}}(s) \approx v_{\pi}(s) & \text{(or } v_*(s)) \\ q_{\mathbf{w}}(s, a) \approx q_{\pi}(s, a) & \text{(or } q_*(s, a)) \end{array}$$

- Update parameter \mathbf{w} (e.g., using MC or TD learning)
- Generalize to unseen states

Agent state

- When the environment state is not fully observable ($S_t^{env} \neq O_t$)
- Use the **agent state**: (with parameters ω)

$$\mathbf{s}_t = u_\omega(\mathbf{s}_{t-1}, A_{t-1}, O_t)$$

- Henceforth, S_t or s_t denotes the agent state
- Think of this as either a vector inside the agent, or, in the simplest case, just the current observation: $S_t = O_t$

Function Classes

Classes of Function Approximation

- **Tabular**: a table with an entry for each MDP state
- **State aggregation**: Partition environment states (or observations) into a discrete set
- **Linear function approximation**
 - Consider fixed agent state update (e.g. $S_t = O_t$)
 - Fixed feature map $x: \mathcal{S} \rightarrow \mathbb{R}^n$
 - Values are linear function of features: $v_w(s) = w^T x(s)$
 - Note: state aggregation and tabular are special cases of linear FA
- **Differentiable function approximation**
 - $v_w(s)$ is a differentiable function of w , could be non-linear
 - E.g., a convolutional neural network that takes pixels as input
 - Another interpretation: features are not fixed, but learnt

Classes of Function Approximation

- In principle, **any** function approximator can be used, but RL has specific properties:
 - Experience is not iid — successive time-steps are correlated
 - Agent's policy affects the data it receives
 - Regression targets can be **non-stationary**
 - ...because of changing policies (which can change the target and the data!)
 - ...because of bootstrapping
 - ...because of non-stationary dynamics (e.g., other learning agents)
 - ...because the world is large (never quite in the same state)

Classes of Function Approximation

- Which function approximation should you choose? This depends on your goals.
 - **Tabular**: good theory but does not scale/generalize
 - **Linear**: reasonably good theory, but requires good features
 - **Non-linear**: less well-understood, but scales well. Flexible, and less reliant on picking good features first (e.g., by hand)
 - (Deep) neural nets often perform quite well, and remain a popular choice

Gradient-based Algorithms

Approximate Values By Stochastic Gradient Descent

- Goal: find w that minimize the difference between $v_w(s)$ and $v_\pi(s)$

$$J(\mathbf{w}) = \mathbb{E}_{S \sim d}[(v_\pi(S) - v_w(S))^2]$$

- Gradient descent:

$$\Delta \mathbf{w} = -\frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w}) = \alpha \mathbb{E}_d (v_\pi(S) - v_w(S)) \nabla_{\mathbf{w}} v_w(S)$$

- Stochastic gradient descent (SGD), sample the gradient:

$$\Delta \mathbf{w} = \alpha (G_t - v_w(S_t)) \nabla_{\mathbf{w}} v_w(S_t)$$

Note: Monte Carlo return G_t is a sample for $v_\pi(s_t)$

Linear function approximation

Feature Vectors

- Represent state by a **feature vector**

$$\mathbf{x}(s) = \begin{pmatrix} x_1(s) \\ \vdots \\ x_n(s) \end{pmatrix}$$

- $x: \mathcal{S} \rightarrow \mathbb{R}^n$ is a fixed mapping from state (e.g., observation) to features
- Short-hand: $x_t = x(S_t)$
- For example:
 - Distance of robot from landmarks
 - Trends in the stock market
 - Piece and pawn configurations in chess

Linear value function approximation

- Approximate value function by a linear combination of features

$$v_{\mathbf{w}}(s) = \mathbf{w}^\top \mathbf{x}(s) = \sum_{j=1}^n \mathbf{x}_j(s) \mathbf{w}_j$$

- Objective function ('loss') is quadratic in \mathbf{w}

$$J(\mathbf{w}) = \mathbb{E}_{S \sim d}[(v_\pi(S) - \mathbf{w}^\top \mathbf{x}(S))^2]$$

- Stochastic gradient descent converges to the **global** optimum
- Update rule is simple

$$\nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t) = \mathbf{x}(S_t) = \mathbf{x}_t \quad \implies \quad \Delta \mathbf{w} = \alpha (v_\pi(S_t) - v_{\mathbf{w}}(S_t)) \mathbf{x}_t$$

Prediction algorithms

- We can't update towards the true value function $v_\pi(s)$
- We substitute a **target** for $v_\pi(s)$
 - For MC, the target is the return G_t

$$\Delta \mathbf{w}_t = \alpha(\mathbf{G}_t - v_{\mathbf{w}}(s)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(s)$$

- For TD, the target is the TD target $R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1})$

$$\Delta \mathbf{w}_t = \alpha(\mathbf{R}_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t)$$

- TD(λ):

$$\Delta \mathbf{w}_t = \alpha(\mathbf{G}_t^\lambda - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t)$$

$$G_t^\lambda = R_{t+1} + \gamma \left((1 - \lambda) v_{\mathbf{w}}(S_{t+1}) + \lambda G_{t+1}^\lambda \right)$$

Monte-Carlo with Value Function Approximation

- The return G_t is an **unbiased** sample of $v_\pi(s)$
- Can therefore apply “supervised learning” to (online) “training data”:
$$\{(S_0, G_0), \dots, (S_t, G_t)\}$$


- For example, using **linear Monte-Carlo policy evaluation**

$$\begin{aligned}\Delta \mathbf{w}_t &= \alpha(\mathbf{G}_t - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t) \\ &= \alpha(G_t - v_{\mathbf{w}}(S_t)) \mathbf{x}_t\end{aligned}$$

- Linear Monte-Carlo evaluation converges to the global optimum
- Even when using non-linear value function approximation it converges (but perhaps to a local optimum)

Control with value-function approximation

Control with Value Function Approximation



Policy evaluation **Approximate** policy evaluation, $q_w \approx q_\pi$

Policy improvement E.g., ϵ -greedy policy improvement

Action-Value Function Approximation

- Approximate the action-value function $q_w(s, a) \approx q_\pi(s, a)$
- For instance, with linear function approximation with **state-action features**

$$q_w(s, a) = \mathbf{x}(s, a)^\top \mathbf{w}$$

- Stochastic gradient descent update

$$\begin{aligned}\Delta \mathbf{w} &= \alpha(q_\pi(s, a) - q_w(s, a)) \nabla_{\mathbf{w}} q_w(s, a) \\ &= \alpha(q_\pi(s, a) - q_w(s, a)) \mathbf{x}(s, a)\end{aligned}$$

Action-Value Function Approximation (Alternative)

- Approximate the action-value function $q_w(s, a) \approx q_\pi(s, a)$
- For instance, with linear function approximation with **state features**

$$\begin{aligned} \mathbf{q}_w(s) &= \mathbf{W}\mathbf{x}(s) & (\mathbf{W} \in \mathbb{R}^{m \times n}, \mathbf{x}(s) \in \mathbb{R}^n \implies \mathbf{q} \in \mathbb{R}^m) \\ q_w(s, a) &= \mathbf{q}_w(s)[a] = \mathbf{x}(s)^\top \mathbf{w}_a & (\text{where } \mathbf{w}_a = \mathbf{W}_a^\top.) \end{aligned}$$

Action-Value Function Approximation

- Should we use action-in, or action-out?
 - Action in: $q_w(s, a) = w^T x(s, a)$
 - Action out: $q_w(s) = Wx(s)$ such that $q_w(s, a) = q_w(s)[a]$
- One reuses the **same weights**, the other the **same features**
- Unclear which is better in general
- If we want to use continuous actions, action-in is easier
- For (small) discrete action spaces, action-out is common (e.g., DQN)

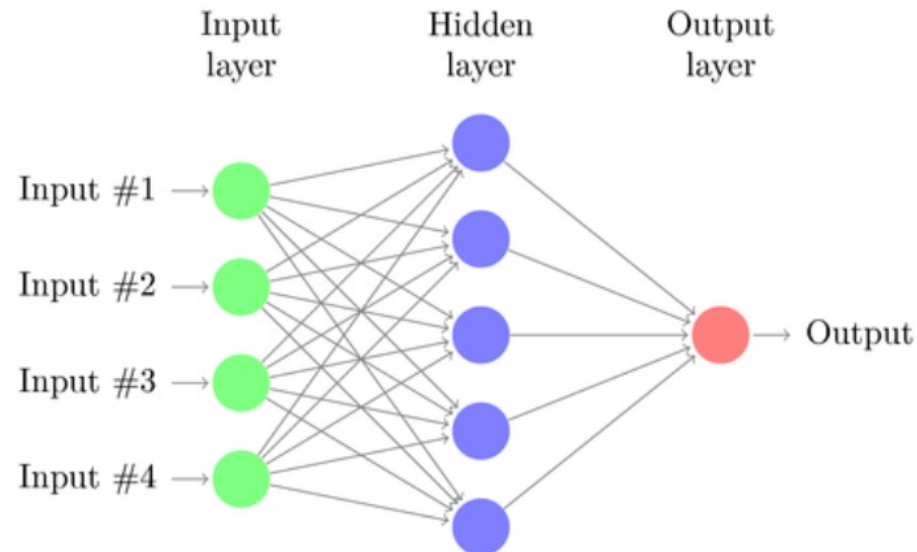
Deep reinforcement learning

RL with Function Approximation

- **Linear value function approximators** assume value function is a weighted combination of a set of features, where each feature a function of the state
- Linear VFA often works well given the right set of features
- But can require carefully hand designing that feature set
- An alternative is to use a much richer function approximation class that is able to directly go from states **without requiring an explicit specification of features**

Neural Networks as Function Approximators

- It is possible to use deep neural networks for function approximations.
- Deep networks are clearly more powerful and can model more complex environments.



Generalization

- Using function approximation to help scale up to making decisions in really large domains



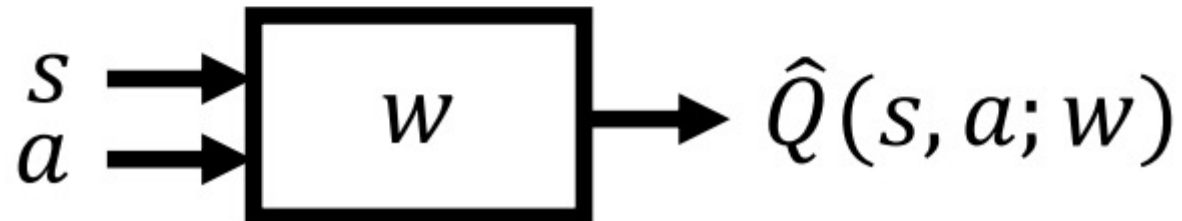
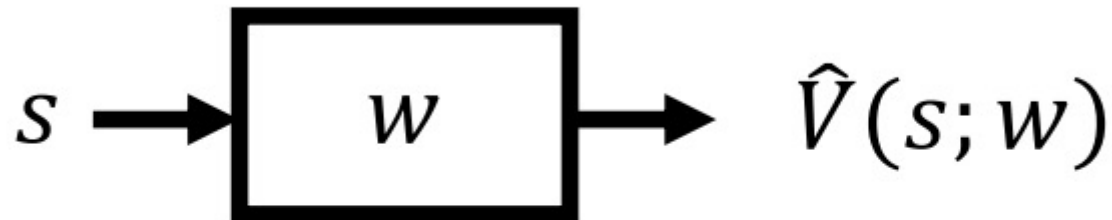
Deep Reinforcement Learning

- Use deep neural networks to represent
 - Value, Q function
 - Policy
 - Model
- Optimize loss function by stochastic gradient descent (SGD)

Deep Q-Networks (DQN)

- Represent state-action value function by Q-network with weights w

$$\hat{Q}(s, a; \mathbf{w}) \approx Q(s, a)$$



Recall: Model free control

- Similar to policy evaluation, true state-action value function for a state is unknown and so substitute a target value
- In Monte Carlo methods, use a return G_t as a substitute target

$$\Delta \mathbf{w} = \alpha(G_t - \hat{Q}(s_t, a_t; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s_t, a_t; \mathbf{w})$$

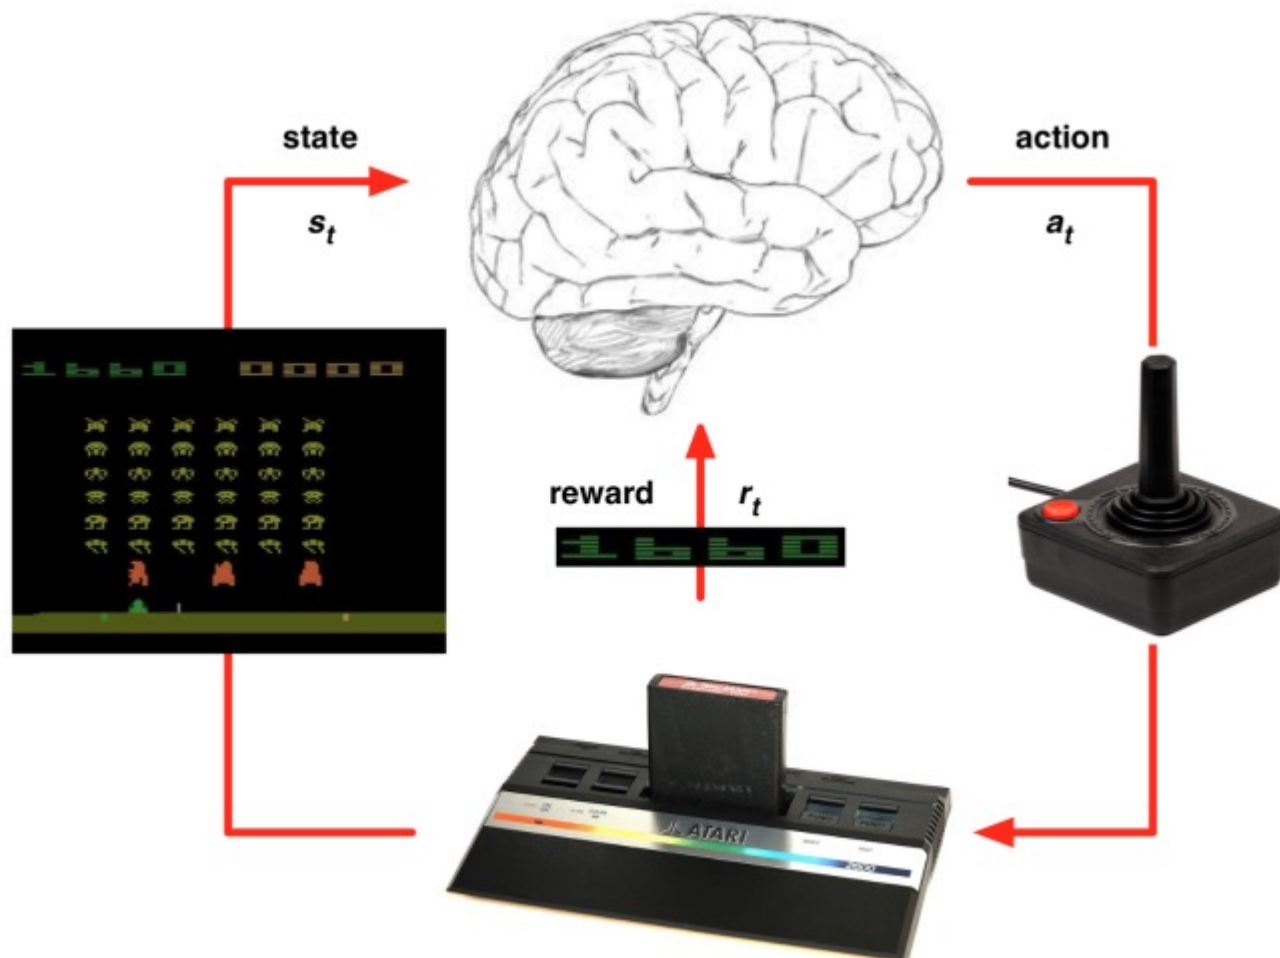
- For SARSA instead use a TD target

$$\Delta \mathbf{w} = \alpha(r + \gamma \hat{Q}(s_{t+1}, a_{t+1}; \mathbf{w}) - \hat{Q}(s_t, a_t; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s_t, a_t; \mathbf{w})$$

- For Q-learning

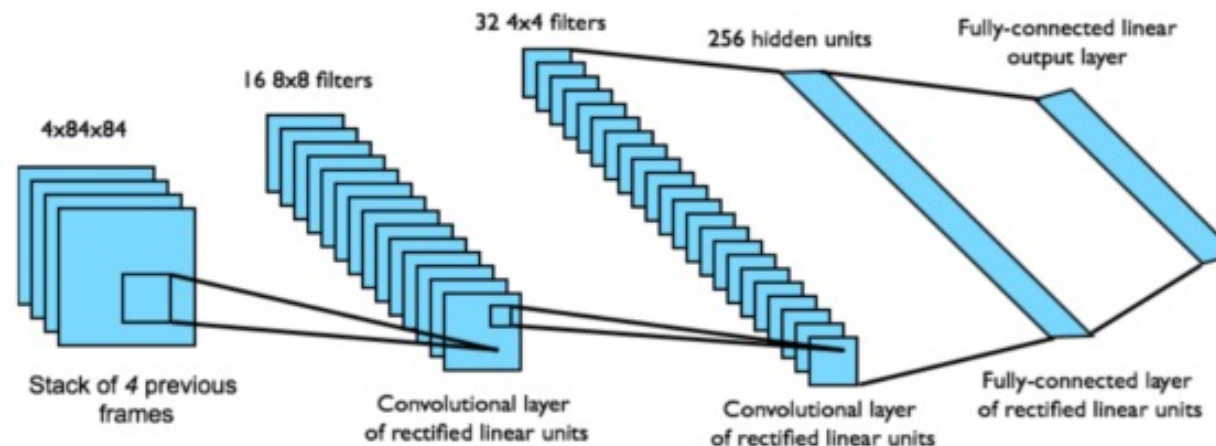
$$\Delta \mathbf{w} = \alpha(r + \gamma \max_a \hat{Q}(s_{t+1}, a; \mathbf{w}) - \hat{Q}(s_t, a_t; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s_t, a_t; \mathbf{w})$$

Doing deep RL in Atari



DQNs in Atari

- End-to-end learning of values $Q(s, a)$ from pixels s
- Input state s is stack of raw pixels from last 4 frames
- Output is $Q(s, a)$ for 18 joystick/button positions
- Reward is change in score for that step
- Network architecture and hyperparameters fixed across all games



DQNs in Atari

- Q-learning converges to the optimal $Q^*(s, a)$ using table lookup representation.
- In value function approximation Q-learning, we can minimize MSE loss by stochastic gradient descent using a target Q estimate instead of true Q (as we saw with linear VFA)
- But Q-learning with VFA can diverge
- Two of the issues causing problems:
 - **Correlations** between samples
 - **Non-stationary** targets
- Deep Q-learning (DQN) addresses these challenges by
 - **Experience replay**
 - **Fixed Q-targets**

DQNs: Experience Replay

- To help remove correlations, store dataset (called a replay buffer) D from prior experience

| | |
|------------------------------|-----------------------------|
| s_1, a_1, r_2, s_2 | \rightarrow s, a, r, s' |
| s_2, a_2, r_3, s_3 | |
| s_3, a_3, r_4, s_4 | |
| ... | |
| $s_t, a_t, r_{t+1}, s_{t+1}$ | |

- To perform experience replay, repeat the following:
 - $(s, a, r, s') \sim D$: sample an experience tuple from the dataset
 - Compute the target value for the sampled s : $r + \gamma \max_{a'} \hat{Q}(s', a'; w)$
 - Use stochastic gradient descent to update the network weights

$$\Delta \mathbf{w} = \alpha (r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

Problem

Can treat the target as a constant scalar, but the weights will get updated on the next round, changing the target value

DQNs: Fixed Q-Targets


- To help improve stability, fix the target weights used in the target calculation for multiple updates
- Target network uses a different set of weights than the weights being updated
- Let parameters w^- be the set of weights used in the target, and w be the weights that are being updated
- Slight change to computation of target value:
 - $(s, a, r, s') \sim D$: sample an experience tuple from the dataset
 - Compute the target value for the sampled s : $r + \gamma \max_{a'} \hat{Q}(s', a'; w^-)$
 - Use stochastic gradient descent to update the network weights

$$\Delta \mathbf{w} = \alpha (r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}^-) - \hat{Q}(s, a; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

DQN Algorithm

1: Input $C, \alpha, D = \{\}$, Initialize $w, w^- = w, t = 0$
 2: Get initial state s_0
 3: loop
 4: Sample action a_t given ϵ -greedy policy for current $\hat{Q}(s_t, a; w)$
 5: Observe reward r_t and next state s_{t+1}
 6: Store transition (s_t, a_t, r_t, s_{t+1}) in replay buffer D
 7: Sample random minibatch of tuples (s_i, a_i, r_i, s_{i+1}) from D
 8: for j in minibatch do
 9: if episode terminated at step $i + 1$ then
 10: $y_i = r_i$
 11: else
 12: $y_i = r_i + \gamma \max_{a'} \hat{Q}(s_{i+1}, a'; w^-)$
 13: end if
 14: Do gradient descent step on $(y_i - \hat{Q}(s_i, a_i; w))^2$ for parameters w : $\Delta w = \alpha (y_i - \hat{Q}(s_i, a_i; w)) \nabla_w \hat{Q}(s_i, a_i; w)$
 15: end for
 16: $t = t + 1$
 17: if $\text{mod}(t, C) == 0$ then
 18: $w^- \leftarrow w$
 19: end if
 20: end loop

linear func.

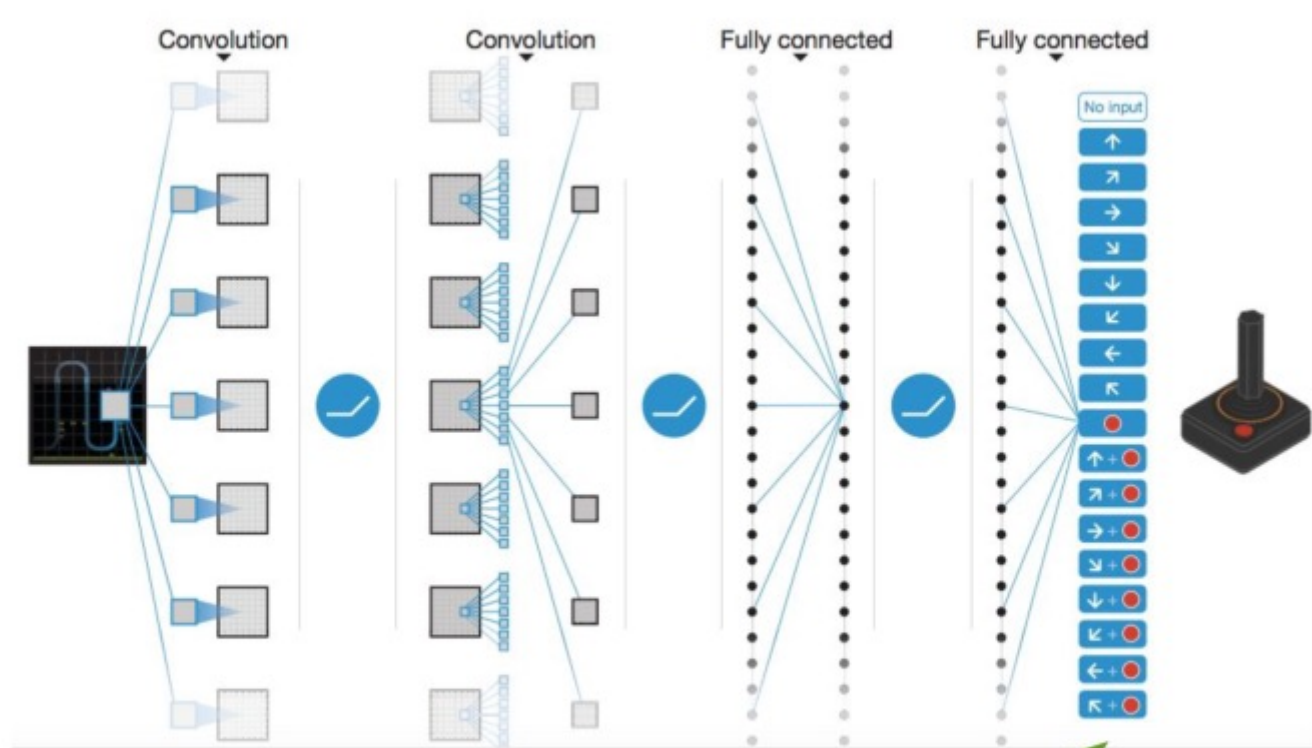


$$\Delta w = \alpha (y_i - \hat{Q}(s_i, a_i; w)) \nabla_w \hat{Q}(s_i, a_i; w)$$

DQN Summary

- DQN uses experience replay and fixed Q-targets
- Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory D
- Sample random mini-batch of transitions (s, a, r, s') from D
- Compute Q-learning targets w.r.t. old, fixed parameters w^-
- Optimizes MSE between Q-network and Q-learning targets
- Uses stochastic gradient descent

DQN



1 network, outputs Q value for each action

Results

| Game | Linear | Deep Network | DQN w/ fixed Q | DQN w/ replay | DQN w/replay and fixed Q |
|----------------|--------|--------------|-------------------|------------------|-----------------------------|
| Breakout | 3 | 3 | 10 | 241 | 317 |
| Enduro | 62 | 29 | 141 | 831 | 1006 |
| River Raid | 2345 | 1453 | 2868 | 4102 | 7447 |
| Seaquest | 656 | 275 | 1003 | 823 | 2894 |
| Space Invaders | 301 | 302 | 373 | 826 | 1089 |